

Capitolis Counterparty Credit Risk Engine

Methods, Results, and How Every Round of Feedback Was Incorporated

Berkeley MFE Industry Project

September 27, 2026

The book, and where every input comes from

The book: 16 trades, 3 counterparties

- 8 equity total-return swaps
- 4 bond forwards, 4 bond total-return swaps
- 37 equities (US + Tokyo) + USDJPY

Every input is real, sourced data

- Equities/dividends: Yahoo Finance
- USD curve: bootstrapped from real Databento SOFR futures
- Treasury yields: FRED
- Long end + vol cube: licensed Bloomberg data, spliced in and cross-checked against public sources

Exposure definition (the brief):

$$\max(V(t+10bd) - V(t-1bd), 0)$$

- Signed prior-day variation margin
- No initial margin, no threshold
- Netted per counterparty only

Secondary view: uncollateralized level exposure $\max(V(t), 0)$ – a no-margin upper bound

The simulation engine, in four bullets

- **Equities & FX:** correlated geometric Brownian motion, real 39×39 correlation matrix (3-year realised)
- **USD rates:** one-factor Hull-White; mean reversion fitted to the swaption market, volatility fitted to realised *long-end* Treasury vol (0.96%, $a=0.0167$)
- **JPY rates:** a second, independent Hull-White factor from the real JPY OIS curve – so JPY names and USDJPY move with real JPY rate risk, not a constant assumption
- **Sampling:** Latin Hypercube throughout; every sensitivity reuses the identical random draws in the base and shocked run (common random numbers)


Scenarios (standard)	5,000
Scenarios (sign-off)	10,000
Reporting horizon	1 year

Automated tests	168
Report length	77 pages


Headline results (5,000 scenarios, close-out exposure)

	MPE99	Peak EE
CPTY_A	\$25.9M	\$4.7M
CPTY_B	\$9.6M	\$1.8M
CPTY_C	\$29.6M	\$5.2M
Portfolio	\$51.0M	\$11.6M

- Portfolio MPE99 = 78% of the sum (\$65.1M): no netting across counterparties, but exposures peak in the same weeks and are positively correlated
- CVA \$11.7k close-out / \$247k level
- SA-CVA capital \$0.10M / \$2.23M
- SA-CCR EAD \$260M




Close-out vs. level exposure: two brackets on the truth



- **Close-out** (the brief, primary): only the 10-day move from an already-margined start. Portfolio MPE99 **\$51.0M**.
- **Level** (secondary): whole mark-to-market at risk, no margin. Portfolio MPE **\$201.9M** – driven by CPTY_C's bond forward.
- The real answer sits between the two, depending on the actual CSA's threshold, MTA, and initial margin – none of which were supplied.

Stress testing methodology




13 fixed scenarios, same random draws as the base case:

- 8 hypothetical: equities $\pm 30\%$, yen $\pm 15\%$, rates $\pm 200\text{bp}$, flight-to-quality, stagflation
- 5 historical replays, chosen by rule (worst 10-day window), not by hand


Read at every reporting date, every trade, every counterparty – not a one-paragraph narrative.

Sensitivities (Greeks): method

- Bump-and-reprice with **common random numbers**: same draws in base and bumped run, so the difference is signal, not noise
- Equity/FX bumps: rescale the simulated GBM paths exactly – **no resimulation**
- Rate/vol bumps: require a genuine resimulation
- Pathwise (closed-form) estimator implemented for equity/FX as a fast cross-check



CVA, DVA, FVA: the credit and funding adjustments



- **CVA:** expected loss from counterparty default, credit-triangle default probabilities from BBB bond-index spread proxies
- **DVA:** mirror image – our own default risk on what we owe them
- **FVA:** funding cost of positive exposure less funding benefit of negative exposure
- **SA-CVA:** Basel standardised-approach capital from CVA's own market-risk sensitivities

Feedback tracker: everything raised, everything addressed

Earlier meeting (5 items)

1. Robust stress testing **Addressed**
2. Sensitivities: method, cost, faster alternative **Addressed**
3. PCA explainability and speed **Addressed**
4. Final results explanation **Addressed**
5. CVA walkthrough and set-up **Addressed**

This meeting (6 items)

6. 2008 and more testing periods **Addressed**
7. LHS to calculate sensitivities **Addressed**
8. KVA addition **Addressed**
9. CVA is for sensitivities, not the number **Addressed**
10. DV01 via a Jacobian **Addressed**
11. Excel parametric vol-bump test **Addressed**

The next 11 slides take each item in turn: what was asked, what we did, the results, and the conclusion.

1. Robust stress testing

Feedback: make it a robust, structured test, not a paragraph – data, inputs, sensitivity along the path, all products

What we did

- 13 scenarios (8 round-number, 5 historical), each a full re-run of the Monte Carlo from the shocked state, same random draws as base
- Results read at every reporting date, every trade, every counterparty
- A per-trade heat map shows exactly which products move under which shocks

Result / finding

- Close-out MPE99 moves at most -24% to $+23\%$; level MPE99 moves up to $+89\%$ – margin is why
- Equity swaps respond only to equity/FX shocks; bond forwards/TRS respond only to rate shocks – confirms trade structure, not a coincidence

Conclusion: stress sensitivity is now quantified, dated, and attributable to a specific shock – not a qualitative claim.

2. Sensitivities: method, cost, faster alternative

Feedback: explain how Greeks are generated, what they cost, and whether something faster exists

Method

- Bump-and-reprice, common random numbers (CRN)
- CRN cuts estimator noise from \$245k to \$1k – 282x

Cost (N=300, 8 cores)

- Base run: 105s
- 78 equity/FX bumps: 88s (vs. ~8,200s naive – 93x cheaper)
- Rate/vol bumps: still need resimulation, ~109s each


Faster alternative

- Pathwise (closed-form) delta for equity/FX: 0.2–0.4s vs. ~260s for bump-and-reprice
- Validated against bump-and-reprice: median 0.02%, worst 0.4% difference

Conclusion: common random numbers is what actually controls Greeks noise; pathwise is a validated, production-ready shortcut for linear (equity/FX) risk only.

3. PCA explainability and speed

Feedback: improve the explainability of the PCA factor model, and quantify the speed it adds



Explainability

- Factor 1 (20% of variance): a broad market factor
- Factor 2 (8.9%): Japan vs. US
- 5 factors explain 47%; 10 factors explain 62%

Speed: measured honestly, found to be none

- Full matrix: 1.91s per correlated-shock step
- 5-factor PCA: 2.08s; 10-factor: 2.29s – *slower*, not faster

Conclusion: kept the full matrix as default; PCA stays as an explainable robustness check, not a speed optimisation.

4. Final results: what drives each counterparty

Feedback: confirm A=equity, C=rates, bonds vs. equities inversely, and why the portfolio is so close to additive

	Equity $\Delta/+1\%$	DV01/ $+1bp$
CPTY_A	-\$2.25M	\$31k
CPTY_B	-\$0.96M	-\$4k
CPTY_C	-\$1.24M	\$594k

Rank correlation of exposures: A-B 0.34, A-C 0.33, B-C 0.12

- A is the most equity-driven; C is the rate netting set (one \$500M short forward on a 2049 Treasury) – but C also carries meaningful equity delta
- Yield-equity correlation in our data is ≈ 0.00 – the “inverse” relationship is regime-dependent, not fixed
- Portfolio MPE99 (\$51.0M) is close to A+C (\$54.4M): nothing nets *across* counterparties by construction, and the three books peak in the same weeks

5. CVA walkthrough and set-up for next session

Feedback: walk through the CVA calculation in full, and set up the open questions for next time

Steps (CPTY_C example)

1. Exposure profile from the simulation, discounted along each path
2. Default probability from a credit spread (BBB proxy: 59bp at 1y, 91bp at 5y) via the credit triangle
3. LGD 60%; multiply and sum


Results

- CPTY_C close-out CVA \$8k; total \$12k
- Rating sensitivity: AA \$7k, A \$8k, BBB \$12k, BB \$18k

Set-up for next time: exposure convention, real ratings/CDS, own credit/funding spread, wrong-way risk, capital approach, margin terms – six explicit open questions carried to today's deck.

6. Add 2008, and more testing periods

Feedback: add maybe 2008, add more testing periods



Method: extended the public price history back to 2007 (28 of 37 names covered) and found the worst 10-day equity window inside two named crises, using the same objective rule as the existing replays.

- **2008 GFC:** Sep 29 – Oct 10, 2008; –23% median / –44% worst single-name return
- **2015 China deval.:** Aug 2015, –6.8% median

Result: 2008 moves close-out MPE99 by –13%, level by +56% – inside the range the existing scenarios already covered. **Conclusion:** leaving out 2008 was not quietly understating the tail.

7. Use Latin Hypercube to calculate a sensitivity

Feedback: is there a way to do the shock as part of this, like calculate sensitivities using this

Method

- In one-factor Hull-White, $r(t) = x(t) + \alpha(t)$
- $x(t)$: the simulated stochastic factor – depends only on the random draws, a, σ
- $\alpha(t)$: purely deterministic, depends only on the curve
- A curve bump (fixed a, σ) leaves $x(t)$ **identical** to the base run


Result

- The whole bump effect is one closed-form number per date, added to the base run's own paths
- **No resimulation. No new random draws.**
- Validated against a true independent resimulation: agreement to 10^{-15} relative precision

Conclusion: yes – exactly, not approximately, for any curve-only bump under the one-factor model.

8. KVA addition

Feedback: KVA addition, CVA, FVA, KVA



Method: added KVA as a fourth, explicitly illustrative xVA term:

$$K(t) = 8\% \times RW \times \alpha \times EPE(t)$$

discounted and integrated like CVA/DVA/FVA, scaled by an assumed cost of capital (8/10/12% swept – none given).

Result at 10% CoC:

- KVA \$22.4k (close-out) vs. CVA \$11.7k
- KVA \$470k (level) vs. CVA \$247k

Conclusion: KVA exceeds CVA because it scales the whole EPE profile by a fixed ratio, not a small IG default probability – explicitly illustrative, pending a real capital methodology.

9. CVA is mostly for understanding sensitivities

Feedback: CVA is done mostly to understand Greeks / sensitivities calculations

What this means in practice

- The headline CVA number is small and not the point
- What a credit/XVA desk actually manages day to day is the *sensitivity* of CVA to market and credit inputs


What we deliver as a result

- CVA01: sensitivity to a 1bp move in the counterparty spread
- Full rating table: AAA \$5.0k → B \$32.8k
- SA-CVA weighted sensitivities by risk class (rates, FX, equity, credit spread)

Conclusion: the report now states this framing explicitly, and the sensitivity chain is the deliverable, not the point estimate.

10. DV01 via a Jacobian: bump the curve, not just the zero rate

Feedback: DV01 rates – use Jacobian, bumping the curve (not just the zero rate?)



Old method: bump the *fitted* zero curve at 8 hand-picked tenors, triangular in between.

New method: bump the curve's own ~45 native construction pillars (SOFR futures + Bloomberg long end), grouped into the same 8 buckets – the genuine par-instrument sensitivity.

Result: the 30y bucket carries **82%** of portfolio DV01 (new) vs. **54%** (old) – the old grid smeared the 20–50y risk (where the 2049 bond forward lives) into the 10y bucket. Bucket sums match the parallel bump to within 0.1% either way.

11. Parametric test in Excel: bump equity vol

Feedback: parametric testing using Excel of bumping vol of equities and showing

Method

- Bump one USD equity's vol $\pm 1\%$
- Compute a simple ATM position's expected exposure $E[\max(S_T - S_0, 0)]$ two ways:
closed-form Black formula (live Excel formula),
and Monte Carlo

Results

- Level match: within $\sim 3\%$
- Vega match: within $\sim 7\%$
- (Realised vol formulas match calibration to 0.1%;
simulated vol matches analytic to 1–2%)

Conclusion: every number in the workbook is a live spreadsheet formula on public data – anyone can reproduce the engine's key numbers with a hand calculator.

Open questions to resolve before the next working session

- Which exposure convention prices CVA in production: close-out, or level – and what does the real CSA specify (threshold, MTA, IM)?
- Can real credit ratings or CDS-implied spreads be obtained for CPTY_A/B/C (the BBB proxy currently drives every credit number)?
- Capitolis' own credit and funding spread for DVA/FVA (currently hypothetical and symmetric)?
- Should wrong-way risk be modelled (exposure/default independence is currently assumed)?
- What real capital methodology and cost of capital should be used for KVA (currently an assumed 8/10/12% grid)?
- Adopt the two-factor G2++ rate model as production default? (Compared: changes portfolio MPE99 by only –3%.)

Thank you

Questions and discussion